

Algorithmic Aspects of Game Theory

Tomasz Garbus

2019

Contents

1	Lecture 1 (27 II 2019)
2	Tutorials 1 (28 II 2019)
3	Lecture 2 (6 III 2019)
4	Tutorials 2 (7 III 2019)
5	Lecture 3 (13 III 2019)
6	Tutorials 3 (14 III 2019)
7	Lecture 4 (20 III 2019)
8	Tutorials 4 (21 III 2019)
9	Lecture 5 (27 III 2019)
10	Tutorials 5 (28 III 2019)
11	Lecture 6 (3 IV 2019)
12	Tutorials 6 (4 IV 2019)
13	Lecture 7 (10 IV 2019)
14	Tutorials 7 (11 IV 2019)
15	Lecture 8 (17 IV 2019)
16	Lecture 9 (24 IV 2019)
17	Tutorials 8 (25 IV 2019)
18	Lecture 10 (8 V 2019)
19	Tutorials 9 (9 IV 2019)
20	Lecture 11 (15 V 2019)

1 Lecture 1 (27 II 2019)

2 Tutorials 1 (28 II 2019)

Hex, choquet: 2 players.

We say a game is **determined** if either of players has a winning strategy. If σ is a winning strategy of P , then $\forall \pi G(\sigma, \pi) \leftarrow$ wins P .

1. Is the choquet game determined if we replace \mathcal{R} with \mathcal{Q} (and its topology)?

If so, who has a winning strategy?

2. Let's consider a variant of choquet games on topological spaces. We have a property: If X is not a Baire space¹ $\implies E$ has a winning strategy (E means Empty, not Eve!).

Example with rational numbers:

$G^q \leftarrow$ set $\mathcal{Q} \setminus q$ dense, open.

Q is countable.

$F \subset Q$

$\cap_{q \in F} G^q = G^F$

$|F| < \mathfrak{c}$

Our strategy:

- we start with set G^{q_0}
- opponent plays a set, say S_1
- we play a set $S_1 \cap G^{q_0} \cap G^{q_1}$

3. If X is complete then NE has w.s.

A complete space is also a Baire space.

4. Consider NIM game.

Setup: n heaps with tokens h_1, h_2, \dots, h_n .

Move: choose a heap and remove $r > 0$ tokens.

Win: The last move.

5 We have two players: E and \forall , Eve move first. Q: Who has a winning strategy? When is the game determined?

$n = 1 \leftarrow$ Eve always wins

$n = 2 \leftarrow ((1, 1) \text{ wins Adam, } (2, 1) \text{ wins Eve, } (2, 2) \text{ wins Adam})$
 $(h_1, h_2) \rightarrow$ equalise them if possible

Eve has a winning strategy iff $h_1 \neq h_2$

General case: Eve wins if the xor of stack heaps is non-zero.

Proof: The winning configuration has xor 0. From a situation with xor $\neq 0$ is always able to produce a situation with xor = 0 and if xor = 0, it's impossible to make a move such that xor = 0 after the move.

1 $(0, \dots, 0, h_j, 0, \dots, 0)$ is a winning position for Eve.

2 if $h_1 \otimes h_2 \dots \otimes h_n = 0$ then the position is *balanced*. Balanced positions are winning positions.

3 Show strategy (next tutorials)

3 Lecture 2 (6 III 2019)

9 Determinacy

If we have a **game of finite duration** with 2 players, we can expand the game in a tree, where a leaf signifies the end of the game. A leaf maps to one of three possible situations:

- existential player (\exists) wins
- universal player (\forall) wins

¹ X is Baire iff:

$G_i \leftarrow$ are dense and open for $i \in \mathcal{N}$ then $\cap_{i>0} G_i \neq \emptyset$

- draw

If we map those situations to values accordingly: 1, -1, 0, the existential player aims to maximize (and universal to minimize) the outcome value.

Let's consider **infinite** games now. Suppose we have 2 players and draw is not possible in the game. If the player does not know the winning strategy, it is possible that they may "lo-op" in a position with winning strategy but never proceed with it.

There **exist** indeterminate perfect information games.

Infinite XOR game: E and A alternately play words $w_0, w_1, w_2, \dots \in \{0, 1\}^+$ which are concatenated to $w_0w_1w_2\dots$

Infinite XOR: any function $f : \{0, 1\}^{\mathbb{N}} \rightarrow \{0, 1\}$ such that if v, w differ by one bit then $f(v) \neq f(w)$.

$v \sim w$ iff differ by a finite number of bits.

We can choose set S s.t. $\{0, 1\}^{\mathbb{N}} \supseteq S$ has $\exists!$ element for each equivalence class (from *Axiom of Choice*).

Each equivalence class of \sim is countable, thus there is continuum of equivalence classes.

Eve wins iff $f(w_0w_1\dots) = 0$, Adam otherwise. No player has a winning strategy in this game.

1. Suppose Adam wins. In the first play:

E	0	w_2	w_4
A	w_1	w_3	

Then in the next game Eve can steal his strategy:

E	1	w_1	w_3	w_5
A		w_2	w_4	

2. Suppose Eve wins. In the first play:

E	w_0	w_2	w_4
A	0	w_3	

Then in the next play:

E	w_0	w_3	
A	1	w_2	w_4

Game on graph

An *arena* is a directed graph, consisting of:

- the set of *positions* Pos
- the set of *moves* $Moves \subseteq Pos \times Pos$

$Pos = Pos_{\exists} \cup Pos_{\forall}$,

$Pos_{\exists} \cup Pos_{\forall} \neq \emptyset$.

A *play* is a finite or infinite sequence of moves:

$q_0 \rightarrow q_1 \rightarrow q_2 \rightarrow \dots \rightarrow q_k (\rightarrow \dots)$.

Game equation

$$\begin{aligned} X &= (E \cap \diamond X) \cup (A \cap \square X) = Eve(X) \\ Y &= (E \cap \square Y) \cup (A \cap \diamond Y) = Adam(Y) \end{aligned}$$

$E = Pos_{\exists}$,

$A = Pos_{\forall}$,

$X, Y \in \mathcal{P}(Pos)$

"Modal logic" symbols here:

$\diamond Z = \{p : (\exists q) Moves(p, q) \wedge q \in Z\}$ ²

$\square Z = \{p : (\forall q)(p \rightarrow q) \Rightarrow q \in Z\}$

Knaster-Tarski Theorem: $\langle L, \leq \rangle$ complete lattice³, $f : L \rightarrow L$ monotonic, then there exists a least fixed point

$\mu x.f(x) = \bigwedge \{d : f(d) \leq d\}$ and a greatest fixed point:

$\nu y.f(y) = \bigvee \{d : d \leq f(d)\}$.

Proof: We show the proof for the greatest fixed point. Let $a = \bigvee A$, $A = \{z : z \leq f(z)\}$

Because f is monotonic, $z \leq a$ implies $f(z) \leq f(a)$. For $z \in A$, this also means $z \leq f(z) \leq f(a)$. Hence, $f(a)$ is an upper bound of A , which follows $a \leq f(a)$. Using the monotonicity of f once more, we obtain $f(a) \leq f(f(a))$. Hence $f(a) \in A$, which follows the converse inequality $f(a) \leq a$.

We consider mappings *Eve* and *Adam* defined in the complete lattice $\langle \mathcal{P}, \leq \rangle$. $Eve(Z)$ is a set of such positions from which Eve can win, and $Adam(Z)$ is a set of such positions from which Adam can win.

Traps and gardens of Eden

A set of positions $Z \subseteq Pos$ is a trap for Adam if $Z \subseteq Eve(Z)$. It means that Adam cannot go out of there.

A set of positions $Z \subseteq Pos$ is Garden of Eden for Eve if $Eve(Z) \subseteq Z$. It means that Adam cannot enter those positions.

The *greatest* trap for Adam is a garden of Eden for Even.

The *least* garden of Eden for Eve is a trap for Adam.

We use the notation: $\bar{Z} = Pos - Z$.

Lemma

$$\overline{Eve(X)} = Adam(\bar{X})$$

Proof. We have:

$$\begin{aligned} \overline{Eve(X)} &= \overline{(E \cap \diamond X) \cup (A \cap \square X)} \\ &= \overline{(E \cap \diamond X)} \cap \overline{(A \cap \square X)} \\ &= (\bar{E} \cup \bar{\diamond X}) \cap (\bar{A} \cup \bar{\square X}) \\ &= (A \cup \bar{\square X}) \cap (E \cup \bar{\diamond X}) \\ &= (A \cap \bar{\diamond X}) \cup (E \cap \bar{\diamond X}) \cup (A \cap E) \cup (\bar{\diamond X} \cap \bar{\square X}) = Adam(\bar{X}) \end{aligned}$$

Proposition: Pos can be divided to three disjoint sets: $\mu X.Eve(X)$, $\mu X.Adam(X)$, $(\nu Y.Eve(Y)) \cap (\nu Y.Adam(Y))$

Definition: strategy

A strategy (for Eve) is a set of finite plays s.t.:

- if $last(w) \in Pos_{\exists}$ then $\exists! q$ s.t. $last(w) \rightarrow q$ and wq is in S

² $p \rightarrow q$ also denotes $Moves(p, q)$ below. A position p , such that $(\forall p)p \not\rightarrow q$ is called *terminal*, which we also write $p \nrightarrow$.

³ A *complete lattice* is a partially ordered set $\langle L, \leq \rangle$, such that each subset $Z \subseteq L$ has the least upper bound $\bigvee Z$, and the greatest lower bound $\bigwedge Z$. In particular, $\bigvee \emptyset$ is the least element denoted \perp , and $\bigwedge \emptyset$ is the greatest element denoted \top .

- if $last(w) \in Pos_{\forall}$ then $\forall(q)(last(w) \rightarrow q) \Rightarrow wq \in S$
- S is closed under initial segments, i.e., if $s_0s_1...s_k \in S$, then $s_0s_1...s_i \in S$, for $0 \leq i \leq k$

4 Tutorials 2 (7 III 2019)

1. Consider NIM game.

We have n stacks of heights $h_1, h_2, \dots, h_n, h \in \{0, 1, 2, \dots\}$.

$W_E = \{x \in \mathbb{N}^R : \text{Eve has a w.s.}\}$.

Let's take $x \in \mathbb{N}^R, x = (h_1, \dots, h_n)$.

$h_1^b \otimes h_2^b \otimes \dots \otimes h_n^b \neq 0 \rightarrow \text{Eve wins from } x$.⁴

The proof consists of 3 observations:

- Final position (all empty stacks) has xor equal to 0.
- From a position s.t. xor $\neq 0$, it is always possible to zero the xor. Let xor be equal some y . Let d be the position of leftmost (most important) bit in binary representation of xor. Thus in some stack, the binary representation must have the d -th bit activated. We can then deactivate d -th bit and set appropriate values on all less significant bits (to zero the xor), and the resulting stack height will be lower.
- From a position s.t. xor = 0, all moves lead to xor $\neq 0$. If we take a non-zero number of tokens from a stack, it means we alter a non-zero number of binary digits in the representation of xor, thus it is no longer 0.

Thus, $W_E = \{x \in \mathbb{N}^R : \otimes x \neq \vec{0}\}$.

2. Represent NIM as a game on graph.

The set of positions is the set of all stack configurations. In general, the space of positions is infinite, but for a fixed game we have finite set of positions.

However, this is not enough (in a game on graph, we want Pos to be disjoint set $Pos = Pos_E \cup Pos_A, Pos_E \cap Pos_A = \emptyset$). Thus we also add information who moves next.

Graph: $G = (V_E \dot{\cup} V_A, E)$.

Strategy is a function $\sigma : V^*V \rightarrow V, \sigma(v_0v_1v_2\dots v_nv_{n+1}) \rightarrow v$.

Let $w = v_1\dots v_n\dots$ be a winning position. We say that first player wins if their strategy $L \in V^w$.

In NIM, the winning condition does not depend on history, so we can collapse the states with the same stacks configuration in the last move (and of course the same current player).

3. Chocolate game.

We have a grid $m \times n$. A player chooses a field and everything on the right and up is erased. One restriction: you cannot choose position $(1, 1)$. The player who makes last move wins.⁵ Is the game determined? Who has a winning strategy?

$1 \times n$: Eve wins (obvious)

$2 \times n$: Eve wins: she eats the position $(n, 2)$ and always maintains the bottom row has one piece more than the top one.

$\omega \times \omega$:⁶ Eve wins: eat the position $(2, 2)$ and then maintain $(1, n); (n, 1)$

⁴ (h^b means binary representation)

⁵ Other definition: players eat chocolate, position $(1, 1)$ is poisoned.

⁶ $\omega = \{0, 1, \dots\}$

$m \times n$:⁷ Assume the 2nd player wins. That means that for any move made in the first move, there exists such move from a second player, that the second player has a winning strategy after the second move.

We can use a strategy stealing argument: First player removes a rectangle larger than 1×1 . Then by assumption the second player has a winning strategy. But instead we can start by removing only one piece of chocolate. Then second player must make a move such that only a single rectangular area is removed (which could be made in one move by the first player). Then first player can copy second player's winning strategy. Thus second player does not have a winning strategy.

Above we have shown that the second player has no w.s.

$\neg \exists \pi \forall \sigma \pi$ wins with $\sigma \Leftrightarrow \forall \pi \exists \sigma \sigma$ wins with π . But this doesn't imply $\exists \sigma \forall \pi \sigma > \pi$ (although there is an implication the other way).

We will just show the game is determined, without showing the winning strategy:

We can build the graph of positions. All paths from $m \times n$ (starting position) are finite, the size of graph is finite as well. We can thus infer the winning position by searching the graph bottom-up (from node $0, 0$ to m, n). We have thus shown the game is determined for a finite size. Since Adam does not have a winning strategy, thus Eve must have it.

For $\omega \times \omega$ we use an argument that after the first move the game graph must have a finite height.

At home: think about chess determinacy, also Armageddon version (black wins if he doesn't lose, also for draw).

5 Lecture 3 (13 III 2019)

Arena: $Pos = Pos_E \dot{\cup} Pos_A$

Move $\in Pos \times Pos$

Zermelo: in chess, either White or Black have winning strategies, or both have a strategy for a draw (at least).

Theorem: In a graph game, for each position γ either Eve has a strategy to win in finite time, or Adam has a strategy to win in finite time, or both have strategies to survive. Moreover, all strategies can be positional⁸.

Strategy (for Eve): *I will tell you what to do if you have obeyed me so far.*

for $w \in S$:

if $last(w) \in Pos_E$ then $(\exists! p)wp \in S$

if $last(w) \in Pos_A$ then $(\forall p)((last(w) \rightarrow p) \Rightarrow wp \in S)$

Strategy S is positional (memory-less) if $last(w) = last(v)$ then $f_S(w) = f_S(v)$.

f_S can be viewed as a partial function on positions of Eve.

If $X \subseteq Eve(X)$ then any function $dom_f = X \cap Pos_E$ such that $(\forall x \in dom_f)f(x) \in X$ witnesses the trap X .

⁷ $m, n \in \mathbb{N}$

⁸ A strategy S is *positional* (or *memory-less*) if this function depends only on the current position.

A function is safe if it is a witness of a trap.

Prop. A partial function $f : Pos_E \supseteq dom_f \rightarrow Pos$ is f_S for some positional strategy iff it is safe. Moreover if f is a witness of trap Z then $f = f_S$ such that $start(S) = Z^9$.

Proof: Suppose f is a witness for trap Z . We define $S = \bigcup_{n=0} S_n$, $S_0 = Z$, S_n

A position p is safe for Eve if $p \in S$ for some strategy for Eve.

Lemma 1 $Safe_E$ – the set of all safe positions for E . $Safe_E = \nu X.Eve(X)$. Moreover, Eve has a positional strategy S , with $start(S) = Safe_E$.

Lemma 2 Win_E^{fin} – the set of positions, from where Eve can win in finite time. Then $Win_E^{fin} = \mu x.Eve(X)$. Moreover Eve has a positional strategy S finitely winning with $Start(S) = Win_E^{fin}$.

Proposition (from Lecture 2) The complement of a trap for Adam is a garden of Eden for him; similarly for Eve.

Proposition (from Lecture 2)

$\mu X.Eve(X) = \nu Y.Adam(Y)$

$\nu X.Eve(X) = \mu Y.Adam(Y)$

Exercise (from Lecture 2) Show that the union of any family of traps for a player is again a trap for this player. Note that this implies that the intersection of any family of gardens of Eden for a player is again a garden of Eden for this player. Which more general property of ordered sets underlines these facts? (Remember the Knaster-Tarski Theorem.)

From definition, a set of positions $Z \subseteq Pos$ is a trap for Adam if $Z \subseteq Eve(Z)$. Let $ZS = Z_1, Z_2, \dots$ be a (not necessarily finite) family of traps. For every Z_i and position $p \in Z_i$, $p \in Eve(ZS)$, thus $ZS \subseteq Eve(ZS)$ so we arrive to the definition of trap again.

6 Tutorials 3 (14 III 2019)

Ordinal numbers:

$0, 1, \dots, \omega, \omega + 1, \dots, \omega + n, \dots, 2\omega, \dots, n\omega, \omega^2, \dots, \omega^5, \dots, \omega^\omega, \dots$

$\alpha \quad \alpha + 1 \mid \alpha \cup \{\alpha\} \quad 0 = \{\}$

Set of ordinals is well founded (no infinite descending sequence).

$\beta \quad \bigcup_{\alpha < \beta} \alpha$

1. NIM on ordinals. $G = (h_1, \dots, h_n), h_i < \omega^\omega$.

1) Is G determined (for which starting positions)?

2) If so, who has a winning strategy?

1. Yes. A graph of states is a DAG and can be divided into disjoint states of winning positions for both players.

2.

a) $h < \omega^\omega$

$h = a_0 + a_1\omega + a_2\omega^2 + \dots + a_n\omega^n$ ($a_i \in \mathbb{N}$) (it is easy to show that this fits under ω^ω)

In finite case ($a_i = 0$ if $i > 0$) we aim to zero the xor of all stacks. We will try to apply this strategy to ordinals. Let's define xor on ordinals.

α, β – ordinals

$\alpha \oplus \beta = (\alpha_0 \oplus \beta_0) + (\alpha_1 \oplus \beta_1)\omega + \dots + (\alpha_n \oplus \beta_n)\omega^n$

Statement In game $G = (h_1, \dots, h_n)$ Eve has a winning strategy iff $\bigoplus_{i>0} h_i \neq \emptyset$. We need to prove that:

1° If Eve plays from $\oplus \neq 0$ then she can always maintain $\oplus = 0$.

2° $\oplus = 0 \Rightarrow$ Every Adam's move makes it $\neq 0$.

3° Ends after a finite number of steps in necessarily Adam's position.

2. Infinite XOR is an undetermined game such that its graph of positions has infinite branching. Try to find a game such that is undetermined but has finite branching.

We modify the game so that player can only place a single letter from the alphabet $\Gamma = \{0, 1, 2, 3, \#\}$ and we have an interpretation function $f : \{0, 1, 2, 3, \#\}^\omega \rightarrow \{0, 1\}^\omega$. If a player does not place a hash (end of a word) in a finite time, they lose. Once a player places a hash, it's the other player's turn.

Gale-Stewart games: $\langle \Gamma, W \subseteq \Gamma^\omega \rangle$, game of perfect information.

Zermelo: Let $G = \langle V, \rightarrow \rangle$ be a graph game. Then there exists a partition $W_E \cup W_A \cup W_N = V$ s.t. player $P \in \{E, A\}$ has a positional winning strategy in position $v \in W_P$.

Reachability game: We select a node in graph, Eve wants to reach that node, Adam does not want to ever reach that node. If the play is infinite and looped without reaching the selected node, Adam wins.

7 Lecture 4 (20 III 2019)

Generalized Zermelo Theorem: For any position p in a deterministic, two-person game with perfect information (and players make moves alternatingly):

- one of the players has a winning strategy (winning in finite time)
- or both players have a surviving strategy

Moreover, those strategies can be positional.

$Eve(X) = (E \cap \diamond X) \cup (A \cap \square X)$ – those positions, from which X can be achieved in one move. If it is Eve's move, there must be one move to position in X , if Adam's – all moves must lead him to X .

We want to find the smallest fixed point $\mu X.Eve(X)$.

$\bigvee f^\zeta(\perp) = \mu x.f(x) = \bigwedge \{d : f(d) \leq d\}$

\perp – minimal element

$f^0(\perp) = \perp$

$f^{\zeta+1}(\perp) = f(f^\zeta(\perp))$

η – limit element:

$f^\eta(\perp) = \bigvee_{\zeta < \eta} f^\zeta(\perp)$

Algorithm for finding winning positions for player $X \in \{Eve, Adam\}$:

Program $Win(X)$ $X \in \{Eve, Adam\}$

W : $Pos \rightarrow \{\perp, \exists, \forall\}$

Forall $p \in Pos$

$W(p) = \perp$

$pred(p) = \emptyset$

$nb(p) = 0$

Forall $(p, q) \in Move$

⁹ $Start(S) = S \cap Pos$

```

pred(q) := pred(q) ∪ {p}
nb(p) := nb(p) + 1
Forall p ∈ PosX̄
  if nb(p) = 0 then Propagate(p, X)

```

```

Propagate(q, X)
  if W(q) = ⊥ then W(q) = X
  Forall p ∈ pred(q)
    if W(p) = ⊥ then
      if p ∈ PosX then
        W(p) = X
        Propagate(p, X)
      else (* p ∈ PosX̄ *) then
        nb(p) := nb(p) - 1
        if nb(p) = 0 then
          Propagate(p, X)

```

Parity games


$Pos_E, Pos_A, Move, rank : Pos \rightarrow C$
 $Win_E \subseteq C^\omega$
 $Win_A \subseteq C^\omega$
 $Win_E \cap Win_A = \emptyset$
 Parity games: $C \subseteq \mathbb{N}$
 $Win_E = \{u \in C^\omega : \limsup_{n \rightarrow \infty} u_n \text{ is even}\}$
 $Win_A = \overline{Win_E}$
 Every position is winning for one of the players, with a positional strategy.

8 Tutorials 4 (21 III 2019)

A game $G = \langle B, W \rangle$
 $B = \langle V, E, \lambda, v_I \rangle$
 $E \subseteq V, \lambda : \rightarrow \Gamma, v_I \in V$
 $W \subseteq \Gamma^* \cup \Gamma^\omega$
 $\Phi : \Gamma^\omega \rightarrow [0, 1]$
 $\Phi_W(x) = \begin{cases} 1, & x \in W \\ 0, & x \notin W \end{cases}$

Example:

Fig. 1: Consider this simple game on graph



$W = \{w \in \{a, b, c\}^\omega \mid \exists_{n < \omega} w = a^n b^n c^\omega\}$ – Adam has a winning strategy (he can infinitely loop in a).

Strategies for Eve and Adam:

$E : \delta : V^* V_E \rightarrow V$ s.t. $\delta(w, v) \rightarrow (v')$ (and $v, v' \in E$)

$A : \pi : V^* V_A \rightarrow V$

Or positional strategies: $\delta : V_E \rightarrow V$

$\pi : V_A \rightarrow V$

A play:

$G < \delta, \pi > \rightarrow v_0 v_1 v_2 \dots v_n \dots = p$

$w = \lambda(v_0) \lambda(v_1) \dots \lambda(v_n) \dots$

$w \in W \Leftrightarrow \delta$ wins with π

(still considering the game above)

$W_2 = W \cup \{a^\omega\}$ – Eve has a winning strategy (but not a positional one)

Reachability game (same graph as before):

$W = \Gamma^* c \Gamma^\omega \cup \Gamma^* c$

$W = V^* F V^\omega \cup V^* F V^*$

$F \subseteq V$ (we want to reach F)

1. Reachability condition, after how many steps every game ends.

In this case, the bound is infinity, since Adam can stay in node a.

2. Create a game that forces the end after 2 steps

(we can extend this for any natural number)

3. Enforce that the game ends after ω steps (you cannot bound it from below).¹⁰



We create such infinite graph:

4. Same for $\omega + 1$ – add one more v_I to the previous graph

5. $\omega + \omega$ – copy the graph, the output from first copy goes into the initial vertex of the second one

$f : V \rightarrow V$

$f(S) = S', S \subseteq S'$

$S' = \{v \in V \mid \forall_{v'} E(v, v') \rightarrow v' \in S\} \cup \{v \in V \mid \exists_{v'} A(v, v') \rightarrow (v, v') \in E\} \cup S$

$f^\alpha(F) = S_\alpha$

$S_0 = F$

$S_\alpha \leftarrow$ set of winning positions of Eve

strenght of game = smallest α s.t. $v_I \in f^\alpha(F)$

6. Given a reachability game $G = \langle V, W \rangle$, compute the set W_E ($W_E = \{v \mid \text{Eve has a winning strategy from } v\}$).

```

S = F;
S' = f(S);
while(S != S') {
  S = S';
  S' = f(S);
}
return S;

```

Parity games

Parity condition:

$\{w \in \Gamma^\omega \mid w = a_0 a_1 a_2 \dots a_n \dots, \limsup a_n \text{ is even}\}$ where $\Gamma \subseteq \mathbb{N}$

Theorem: Parity games are positionally determined.

¹⁰ The smallest number of steps required for reachability game to finish is sometimes called the **strength of a reachability game**.

7. Let G be a parity game, i.e. $G = \langle B, \text{PARITY} \rangle$, s.t. $V_A = \emptyset$ (only Eve moves). Compute the set of winning positions for Eve.¹¹

Observation $v \in V$ is a winning position iff we can reach from v a loop $x_0 \dots x_n$ s.t. the highest priority in the loop is even. Proof

\Leftarrow easy

$\Rightarrow \exists \rightarrow (v = v_1)v_2v_3\dots v_n\dots$ with ranks $a_1a_2a_3\dots a_n\dots \rightarrow$ winning. The highest priority occurring infinitely often is even, let's call it a .

Thus $\exists : \forall_{n_0} \forall_{n \geq n_0} a_n \leq a$


Take a sequence $a_{n \geq n_0} \dots a \dots a \dots a \dots a$. So we choose a path between two occurrences of a and loop it.

HOME Describe the algorithm basing on the observation.

9 Lecture 5 (27 III 2019)

Parity games: Eve wins a play $\pi = p_0p_1\dots$ if $\limsup_{n \rightarrow \infty}(\text{rank}(p_n))$ is even, Adam otherwise.

Fig. 2: What are W_E, W_A here? (squares are Adam, circles Eve)



General games on labelled graphs:

$Pos_E, Pos_A, Move \in Pos \times Pos$


$rank : Pos \rightarrow C$

Winning criteria: $W_E \subseteq C^\omega, W_A \subseteq C^\omega, W_E \cap W_A = \emptyset$

Theorem: Parity games are positionally determined. We prove for finite arenas. Some positions are marked \top immediately winning for Eve, \perp immediately winning for Adam. Induction on $\#$ positions, not marked \top, \perp .

Induction step: Choose a position p with the highest possible rank d , assume (wlog) d is even.

- 1 Suppose $p \in Pos_E$. Mark p by \top .



- (1a) There is a move from p to W_E .
- (1b) All moves from p go to W_A .

- 2 $p \in Pos_A$

- (2a) All moves from p go to W_E .
- (2b) There is a move from p to W_A .

Parity game – n vertices, d ranks. "Simple" algorithm – $n^{\frac{d}{2} + \Theta(1)}$

¹¹Since we are looking for an algorithm, assume graph G is finite in size.

10 Tutorials 5 (28 III 2019)

1. Problem

In: a parity game G and a positional strategy of Eve σ
Out: is σ winning?

σ is winning $\leftrightarrow \forall \pi G(\sigma, \pi) \in \text{PARITY}$

$\sigma : V \rightarrow V$

2. In: a parity game G , and an initial position v_I .

Out: Is v_I winning for Eve?

Solution in NP:

- guess the positional strategy of Eve $\rightarrow \sigma_p$
- check if σ_p is winning

Solution in coNP – check for strategy for Adam.

The problem is (at least) in $\text{NP} \cap \text{coNP}$

3.

Muller $C \subseteq \mathbb{N}, \mathcal{F} \subseteq 2^C$

$C^\omega \ni p = a_1a_2\dots a_n\dots$ if $\text{Inf}(p) \in \mathcal{F} \leftarrow$ set of letters in p that appears ∞ -often


Parity – $\max(\text{Inf}(p))$ is even

- 1) Parity is Muller
- 2) Muller is not positionally determined
- 3) Muller is not Parity
- 4) Show that Muller is determined, what is the required memory? Show that if G is a Muller game, then for every initial vertex v_I , one of the players has a winning strategy.

Solutions:

- 1) Parity game of index (i, k) can be represented as a Muller game, where \mathcal{F} consists of subsets of $\{i, i+1, \dots, k\}$ with even maximum.

- 2) Example below:



$\mathcal{F} = \{\{a, b, c, d\}\}$

Every positional strategy removes either b or c and we require both.

- 3) From (2)

- 4) Idea: We will try to transform it to some Parity game, i.e.

$G = \langle V, E, \lambda, v_I, \mathcal{F} \rangle \rightsquigarrow G' = \langle V', E', \lambda', v'_I, \text{PARITY} \rangle$

$E \subseteq V \times V$

$\lambda : V \rightarrow C$

s.t. v_I is winning for Eve in $G \Leftrightarrow v'_I$ is winning for Eve in G' (we need the same for Adam). In fact, we only need the implication to the left side.

LARs

LAR (latest appearance record) $w \in C \cup \{\natural\} (= \Gamma)$, every letter occurs at most once, \natural occurs always.

up : $\text{LAR} \times \Gamma \rightarrow \Gamma$

$$\text{up}(v_i \natural w_i, a) = \begin{cases} v_i \natural w_i a & a \notin v_i, a \notin w_i \\ [v_i w_i]_{a \mapsto \natural, \natural \mapsto \emptyset} a & \text{else} \end{cases}$$

Theorem: The following cylindrification of G is a solution to task

$$(4): \text{rank}(p, v \natural a_1 \dots a_l) = \begin{cases} 2l & \text{if } \{a_1, \dots, a_l\} \in \mathcal{F} \\ 2l + 1 & \text{otherwise} \end{cases}$$

11 Lecture 6 (3 IV 2019)

G^{++}

$\{0, 1, \dots, n\}^{k+1} \approx k + 1$ -digit numbers in base $n + 1$

Overflow = $(n + 1)^{k+1}$

$$\text{up}(2i + 1, \underbrace{a_0 a_1 a_2 \dots a_k}_m) = m + (n + 1)^i$$

$$\text{up}(2i, a_0 a_1 \dots a_k) = 0 \dots 0 a_i a_{i+1} \dots a_k$$

G^{++} is equivalent to G (i.e. if Adam has winning strategy in G , he has a winning strategy in G^{++})

12 Tutorials 6 (4 IV 2019)

Last week

Theorem If G is a finite Muller game, then one of the players has a finite memory winning strategy. $\langle V, E, v_I, \text{MULLER} \rangle$

Proof G' – a parity game s.t.

Adam wins in $G' \Rightarrow$ Adam wins in G

Eve wins in $G' \Rightarrow$ Eve wins in G

LARs

Automata

$\mathcal{A} = \langle \Gamma, Q, q_i, \delta, F \rangle$

Word: $q_I \in Q, \delta : Q \times \Gamma \times Q, (q, a) \rightsquigarrow \{q_1, q_2, q_3, \dots\}$

Infinite automata

Büchi condition $F \subseteq Q$ (the states appearing infinitely often)

Müller condition $F \subseteq 2^Q$

Parity

Safety – Parity with condition $[0]$

Theorem Every regular language of infinite words is recognizable by a deterministic Müller automaton.

1. Show that if G is a finite Müller game, then set of plays is regular.

P – set of plays

$P \subseteq V^\omega$

The automaton is the graph. $\langle V, V, v_I, \{(v_1, v_2, v_3) \mid E(v_1, v_2)\} \rangle$

Def Game is ω -regular if the winning condition is a regular set of winning plays.

$G = \langle V, L \rangle. (L \subseteq \Gamma^\omega, \text{Eve wins if } p \in L)$

2. Show that if G is ω -regular finite game then one of the players has a finite winning strategy.

$G = \langle E, V_I, V_A, V_E \rangle, V = V_A \cup V_E$

$G_A = \langle V'_A = V_A \times Q, V'_E = V_E \times Q, V'_I, E' \rangle$

$G_A = \langle V', E', V'_I \rangle$

$V' = V \times Q$

$E' = \{((v_1, q_1), (v_2, q_2)) : (v_1, v_2) \in E, q_2 = \delta(q_1, v_1)\}$


$v'_I = (v_I, q_I)$

$W = \text{inf}(L(P)) \in \mathcal{F}$

$L = \pi_2(p)$

Lemma: E has a finite memory winning strategy in $G_A \Rightarrow E$ has a winning strategy in G .

13 Lecture 7 (10 IV 2019)



Parity games $n, \underbrace{d}_{\text{rank}}$.

Karoliina Lehtinen – algorithm.

G – a finite parity game.

We create a game $\mathcal{R}_k^E(G)$ with k registers containing priorities

$\text{Mem} \subseteq \{0, 1, \dots, d\}^k. \text{Pos}_E = (\text{Pos} \times \text{Mem} \times \{0\}) \cup (\text{Pos}_E \times \text{Mem} \times \{1\})$

$\widetilde{\text{Pos}}_A = \text{Pos}_A \times \text{Mem} \times \{1\}.$

If $p \rightarrow q$ in G then

$$(p, \alpha, 1) \xrightarrow{1} (q, \text{up}(\alpha), 0)$$

$$(p, \alpha, 0) \rightarrow (p, \text{reset}(\alpha), 1)$$

$$(p, \alpha, 0) \xrightarrow[1]{\text{skip}} (p, \alpha, 1)$$

$$[\text{up}(\alpha)]_i = \max(\alpha_i, \text{rank}(q))$$

Lemma 1 If Adam wins original game G from position p , he wins $\mathcal{R}_k^E(G)$ from position $(p, \alpha, 1)$ for any number of registers k . Adam plays the strategy from original game, we will show that any play is winning. Let q be the maximal odd rank. Let i be the deepest register, which resets infinitely often. We will show, that infinitely often, in the moment of reset of the register i , it contains q .

Lemma 2 If Eve wins the game G from position p , then (\exists_k) Eve wins $\mathcal{R}_k^E(G)$ from position $(p, \alpha, 0)$.

k : number of ranks from the original game.

To each even rank d we dedicate some register. When in original game there appears an odd rank d , Eve resets the register dedicated to rank d .

14 Tutorials 7 (11 IV 2019)

Büchi-Landweber thm.

Let G be an ω -regular game, i.e. the winning strategy is ω -regular. Then, either of the players has a winning strategy.

$\langle V, E, v_I, \lambda, L \rangle$
 $\lambda(v_0 \dots v_n \dots) \in L$.

- 1) L is regular $\Rightarrow L$ recognized by Deterministic Müller Automaton \mathcal{A} .
- 2) G' is a Müller Game created as a product of $(V, E, \lambda) \times \mathcal{A}$.
- 3) Adam/Eve wins in $G' \Rightarrow$ Adam/Eve wins in G .

Complexities

Problem 1.


In. $G = \langle B, L(A) \rangle \leftarrow \omega$ -regular game.

Out. Does Eve win?

- a) A is a deterministic parity automaton.
 The upper bound is $NP \cap coNP$, because G is a parity game.
 Algorithm: Reduce G to $G' = G \times A$ parity game, solve G' .
 Lower bound: can encode parity games.
- b) A is a DMA.
- c) A is a non-deterministic Büchi automaton.
 Algorithm: compute DPA A' such that $L(A) = L(A')$, $|a|^{|a|}$, $[0, 2|a|]$, use (a). This is in EXP.

Flow Games

Fig. 3: Example flow game



$G = (V, E, E_\diamond \dot{\cup} E_\square, \lambda, e \in E)$.

$\lambda : E \cup V \rightarrow [0, d]$

Valid is $\forall_v \lambda(v) \leq \sum_{(a,v) \in E} \lambda(a, v)$.

A player $P \in \{\diamond, \square\}$ chooses an edge $e \in E_p$, and flips it.

Player \diamond wins if a predefined edge e is flipped.

Problem 2. Show that if there is only one player \diamond , and every edge can be flipped at most once, then the game is NP-complete.


In NP: It is in NP because we can guess the witness in NP and verify in polynomial time.

Complete: Reduction 3-SAT¹² \rightarrow Game.

Problem 3. Show that if we allow that an edge is flipped at most once, then the game is APTIME-complete, PSPACE-complete.

¹² 3-SAT: Formula $F_1 \wedge F_2 \dots \wedge F_n$, s.t. $F_i = x_{i_1} \vee x_{i_2} \vee x_{i_3}$ s.t. $x_j = y_i$ (or $\neg y_i$) for $y_1 \dots y_k$.

15 Lecture 8 (17 IV 2019)



Karoliina Lehtinen algorithm

16 Lecture 9 (24 IV 2019)

Mean payoff games

Adam and Eve play on a weighted graph (negative weights also occur). The play ends once one vertex has been visited twice (there has occurred a loop). The score of the game is the mean of the weights **only** on the loop. Eve wants to maximize the score, Adam - to minimize. The game is of course determined.

Arena: finite graph $\langle Pos_E, Pos_A, Move \subseteq Pos \times Pos \rangle$

$c : Move \rightarrow \mathbb{Z}$

Suppose an infinite play $\pi = (p_0, p_1, p_2, \dots)$ is taken.

Then:

$$Pay_E(\pi) = \liminf_{n \rightarrow \infty} \frac{1}{n} \sum_{i=0}^{n-1} c(p_i, p_{i+1})$$

$$Pay_A(\pi) = \limsup_{n \rightarrow \infty} \frac{1}{n} \sum_{i=0}^{n-1} c(p_i, p_{i+1})$$

$Pay_A(\pi) - Pay_E(\pi)$ is taken by Tax Office.

Γ - mean-payoff game

G - finite game (terminates once the first loop is closed and the value is the mean on this loop)

Thm (1) G is determined (obvious), for each position p there is an optimal game value $val_G(p)$, such that Eve and Adam can both achieve precisely this.

Thm (2) Γ is determined and $val_G(p) = val_\Gamma(p)$.

Eve has a strategy S in G . We define a strategy S' in Γ . Play S , as soon as a loop is formed, remove it and play S again.

Claim: with this strategy Eve can achieve at least $val_G(p)$.

Note:

$$\frac{a_1 + \dots + a_m}{m} \geq c, \frac{b_1 + \dots + b_n}{n} \geq c, \text{ then: } \frac{a_1 + \dots + a_m + b_1 + \dots + b_n}{m+n} \geq c$$

Therefore:

$$\frac{1}{n} \sum_{i=0}^{n-1} c(p_i, p_{i+1}) = \frac{\text{no loop}}{n} + \text{mean loop} + \text{mean loop} + \dots$$

Lemma (orange)

q_0, p - positions

$G_{q_0}(p)$ - Play as in G , but if you reach q_0 , reset¹³.

$val_{G_{q_0}}(p) = val_G(p)$

Strategy for Eve in G_{q_0} from p : Play S (the strategy in G) from p , but if we visit q_0 , play S from q_0 .

¹³ When players eat the orange, they forget about the past :D

17 Tutorials 8 (25 IV 2019)

AFA (alternating finite automaton) $\mathcal{A} = \langle \Gamma, Q, q_I, \lambda, \delta, \text{Win} \rangle$

Γ – finite

$q_I \in Q$

$\lambda : Q \rightarrow [m, n]$

$\delta : Q \times \Gamma \rightarrow B^+(Q \times \{+1\})$

$w = \Gamma^\omega$

$G(\mathcal{A}, w) :$

$V = B^+(Q \times \{+1\}) \times \text{pos}(\omega)$

$v_I = (q_0, 0)$

$E :$

Eve:

$((q+1), n) \rightarrow (\delta(q, w[n]), n+1)$

Eve:

$(\phi \vee \psi, n) \rightarrow (\phi, n)$

$(\phi \vee \psi, n) \rightarrow (\psi, n)$

Adam:

$(\phi \wedge \psi, n) \rightarrow (\phi, n)$

$(\phi \wedge \psi, n) \rightarrow (\psi, n)$

Eve:

$(q_0, 0) \rightarrow (\delta(q_0, w[n]), 1)$

A accepts w if Eve wins in $G(A, \omega)$

1. Show that ATA are closed under \cup, \cap, \neg .

Nonemptiness of an AFA is PSPACE-complete.

Hint for homework: For EXPTIME-complete, you can use the following "generic" problem:

In: \mathcal{M} – a (D/N/A) Turing machine, $n \in \mathbb{N}$ in unary

Out: Does \mathcal{M} accept an empty input using n steps, n memory cell.

For deterministic machine it is PTIME-complete (for steps and PSPACE for memory), for nondeterministic it is (NP-complete for steps and NPSPACE=PSPACE-complete for cells). For alternating machine it is APSPACE=EXP.

Mean-payoff games

$G = \langle V, E, \lambda, v_I \rangle$

$\lambda : E \rightarrow \mathbb{Z}$

$p = v_0 v_1 \dots v_n \dots$


PG: $\lambda(p \in W)$

$\Phi : \text{plays}(G) \rightarrow [0, 1]$

$\lim_{n \rightarrow \infty} \frac{\sum_{i=0}^n \lambda(e_i)}{n}$

	MPG	PG
	$\Phi(p)$	$\text{Parity} \ni P$
δ	optimizing strategy $\neg \exists$	winning strategy \exists

$\Phi_{\text{Parity}}(p) = \begin{cases} 1 & \text{parity is even} \\ 0 & \text{if not} \end{cases}$



$$\underline{val} = \sup_{\sigma} \inf_{\pi} G(\sigma, \pi), \quad \overline{val} = \inf_{\pi} \sup_{\sigma} G(\sigma, \pi)$$

a) $\overline{val} \geq \underline{val}$, if $=$ then G is determined

b) If $\text{Im}(\phi) = \{0, 1\}$ then there is an optimal strategy for one of the players (or a winning strategy). (Assume $\overline{val} = \underline{val}$)

18 Lecture 10 (8 V 2019)

Mean payoff games

We know: for each position p , there is a compromise value

$\text{val}_G^E(p) = \text{val}_G^E(p)$ (determinacy)

Moreover, a strategy in G can be used for Γ .

Challenge: optimal strategies can be positional.

Orange game

G_{q_0}

Restart in q_0 .

Lemma: $\text{val}_G(p) = \text{val}_{G_{q_0}}(p)$

Observation: If optimal strategy S from point p in game G "enters" point q , then $\text{val}_\Gamma(p) \leq \text{val}_\Gamma(q)$.

19 Tutorials 9 (9 IV 2019)

20 Lecture 11 (15 V 2019)

Topic: Nash equilibrium

Fig. 4: A sample payoff table for 2 players and 2 strategies

	B	S
B	6	1
S	3	5

$$\sup_x \inf_y f(x, y) \stackrel{?}{\leq} \inf_y \sup_x f(x, y)$$

$$\inf_y f(x_0, y) \leq f(x_0, y_0) \leq \sup_x f(x, y_0)$$

$$\sup_x \inf_y f(x, y) \leq \inf_y \sup_x f(x, y)$$

S_E, S_A – sets of strategies for Eve and Adam

$$f(s_E, s_A) = \begin{cases} 1, & \text{Eve wins} \\ 0, & \text{Adam wins} \end{cases}$$

$$\sup_{s_E} \inf_{s_A} f(s_E, s_A) \stackrel{?}{=} \inf_{s_A} \sup_{s_E} f(s_E, s_A)$$

The above equality holds if and only if the game is **determined**.

Nondeterminacy:

$$\forall_{s_E} \exists_{s_A} f(s_E, s_A) = 0$$

$$\forall_{s_A} \exists_{s_E} f(s_E, s_A) = 1$$

Then:

$$\inf_{s_A} \sup_{s_E} f(s_E, s_A) = 1$$

$$\sup_{s_E} \inf_{s_A} f(s_E, s_A) = 0$$

Consider the following game:

There is a scholarship to be granted to a single student. Let S be the set of students, P be the set of classes. Each student's value is $\Theta(p, s) \cdot 100$ PLN, where $\Theta(p, s)$ is the student's grade from class p . So the value of scholarship depends on the student's grade.

The student council chooses the student (s) and the financial dean chooses the class (p).

Consider variant, where both sides make their choices independently. When will neither side regret their choice?

If (and only if):

$$\sup_s \inf_p \Theta(s, p) = \inf_p \sup_s \Theta(s, p)$$


Another example: pollution game:

There are N countries. A country which controls air pollution, pays 3. A country which does not control air pollution, pays 1 to everyone.


Without control, everyone pays $N - 1$. With control, everyone pays 3.

Braess paradox

There are two roads to get from A to B . There are $N = 100$ people driving. The time of traversing the edge with 1 is 1 hour. The time of traversing the edge with x is $\frac{x}{N}$ hours, where x is the number of people who have chosen this edge.



The equilibrium exists – drivers divide in two equal groups. Consider what happens, when a superfast edge is added:



Tragedy of commons

N people want to appear on television.

Fig. 5: A replica of prof. Niwiński's drawing



x_i is a time i -th player speaks. If $\sum x_i = 1$, the player does not have time.

If $\sum x_i < 1$, then time i is $x_i(1 - \sum_{j=1}^N x_j)$.

i -th reasoning:

If $\sum_{j \neq i} x_j = t$, then maximize $-x^2 - (t - 1)x$. He or she chooses $x = \frac{1-t}{2}$.

We get a set of equations:

$$2x_1 + x_2 + x_3 + \dots + x_N = 1$$

$$x_1 + 2x_2 + x_3 + \dots + x_N = 1$$

...

$$x_1 + x_2 + x_3 + \dots + 2x_N = 1$$

Then $x_i = \frac{1}{N+1}$. i -th time live is $\frac{1}{(N+1)^2}$. Total time is $\frac{1}{N+1}$. If everyone chose a smaller value, $\frac{1}{2N}$, then total time would be $\frac{1}{2}$.